



Machine Learning with Scikit-Learn

scikit-learn

Simple and efficient tools for data mining and data analysis

Accessible to everybody, and reusable in various contexts.

Built on **NumPy, SciPy, and matplotlib**


Representation of Data in Scikit-learn

Machine Learning is all about creating models from data.

`[n_sample, n_feature]`

<https://archive.ics.uci.edu/ml/datasets/iris>

Nearest Neighbors Classification



sklearn.neighbors.KNeighborsClassifier

n_neighbors :int, optional (default = 5)

Number of neighbors to use

weights : str or callable, optional (default = 'uniform')

weight function used in prediction. Possible values:

'uniform' : uniform weights. All points in each neighborhood are weighted equally.

'distance' : weight points by the inverse of their distance.

fit(X, y)

Fit the model using X as training data and y as target values


predict(X)

Predict the class labels for the provided data

Iris flower data set

The data set consists of 50 samples from each of three species of Iris (*Iris setosa*, *Iris virginica* and *Iris versicolor*).

Classes	3
Samples per class	50
Samples total	150
Dimensionality	4
Features	real, positive




Iris Versicolor



Iris Setosa



Iris Virginica



Dataset Order ⇅	Sepal length ⇅	Sepal width ⇅	Petal length ⇅	Petal width ⇅	Species ⇅
1	5.1	3.5	1.4	0.2	<i>I. setosa</i>
2	4.9	3.0	1.4	0.2	<i>I. setosa</i>
3	4.7	3.2	1.3	0.2	<i>I. setosa</i>
4	4.6	3.1	1.5	0.2	<i>I. setosa</i>
5	5.0	3.6	1.4	0.3	<i>I. setosa</i>
6	5.4	3.9	1.7	0.4	<i>I. setosa</i>
7	4.6	3.4	1.4	0.3	<i>I. setosa</i>
8	5.0	3.4	1.5	0.2	<i>I. setosa</i>
9	4.4	2.9	1.4	0.2	<i>I. setosa</i>
118	7.7	3.8	6.7	2.2	<i>I. virginica</i>
119	7.7	2.6	6.9	2.3	<i>I. virginica</i>
120	6.0	2.2	5.0	1.5	<i>I. virginica</i>
121	6.9	3.2	5.7	2.3	<i>I. virginica</i>
122	5.6	2.8	4.9	2.0	<i>I. virginica</i>
123	7.7	2.8	6.7	2.0	<i>I. virginica</i>

```
from sklearn.datasets import load_iris
iris=load_iris()
print(type(iris))
print(iris.data)
print(iris.data.shape)
print(iris.feature_names)
print(iris.target_names)
print(type(iris.data))
print(type(iris.target))
```

Each row is an Observation

(also know as: sample , example ,instance , record)

Each Column is a feature

(also Know as: predictor , attribute , independent variable)

```
from sklearn.datasets import load_iris
iris=load_iris()
iris.keys()
iris['data']
iris['feature_names']
iris['data'].shape
print(iris["target"])
print(iris["target_names"])
```

Predicting the Iris Flower

1.Import the class

```
from sklearn.datasets import load_iris  
from sklearn.neighbors import KNeighborsClassifier
```

2.Instantiate an Estimator

```
knn=KNeighborsClassifier(n_neighbors=1)
```

3.Model Training

Learning is the Relationship Between Data and Target

```
iris=load_iris()
```

```
knn.fit(iris.data,iris.target)
```

Predict The Response for New Observation

The new observation are called "out-of-sample" Data

```
knn.predict([[3,5,4,2],])  
print(iris.target_names[_])
```

```
iris_new=[[3,5,4,2],[5,4,3,2]]  
knn.predict(iris_new)  
print(iris.target_names[_])
```

Regression

The data will be split into a training and test set. Once we have the test data, we can find a **best fit line** and **make predictions**.

Load Data

Split the data into training/testing sets

Split the targets into training/testing sets

Create linear regression object

Train the model using the training sets

```
import matplotlib
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model
import pandas as pd
```

```
df =
pd.read_csv(r"C:\Users\Shubh_Ram\Desktop\Workshop_
AI_CODE\data\Housing _Data.csv")
```

df

```
Y = df['price']
X = df['lotsize']
```



```
X=X.reshape(len(X),1)
Y=Y.reshape(len(Y),1)
```

```
# Split the data into training/testing sets
```

```
X_train = X[:-250]
```


```
X_test = X[-250:]
```

```
# Split the targets into training/testing sets
```

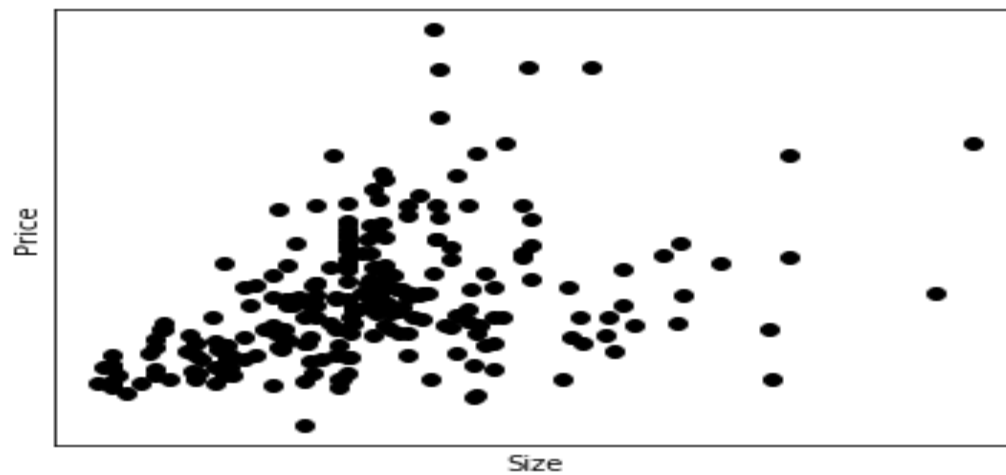
```
Y_train = Y[:-250]
```

```
Y_test = Y[-250:]
```


```
plt.scatter(X_test, Y_test, color='black')  
plt.title('Test Data')  
plt.xlabel('Size')  
plt.ylabel('Price')  
plt.show()
```



Test Data



Test Data




```
regr = linear_model.LinearRegression()  
regr.fit(X_train, Y_train)  
plt.plot(X_test, regr.predict(X_test),  
color='red',linewidth=3)  
plt.show()
```

```
print( str((regr.predict(5000))) )
```

```
# Create linear regression object
regr = linear_model.LinearRegression()
# Train the model using the training sets
regr.fit(X_train, Y_train)
# Plot outputs
plt.scatter(X_test, Y_test, color='black')
plt.title('Test Data')
plt.xlabel('Size')
plt.ylabel('Price')
plt.xticks(())
plt.yticks(())
plt.plot(X_test, regr.predict(X_test), color='red',linewidth=3)
```

The Digit Dataset

This dataset is made up of 1797 8x8 images. Each image, like the one shown below, is of a hand-written digit. In order to utilize an 8x8 figure like this, we'd have to first transform it into a feature vector with length 64.



```
from sklearn import datasets
digits=datasets.load_digits()
x=digits.data
y=digits.target
print(x)
print(y)
```

Neural network models (supervised

Class `MLPClassifier` implements a multi-layer perceptron (MLP) algorithm that trains using Backpropagation.

MLP trains on two arrays: array `X` of size `(n_samples, n_features)`, which holds the training samples represented as floating point feature vectors; and array `y` of size `(n_samples,)`, which holds the target values (class labels) for the training samples

MLPClassifier

1.Hidden Layer

2.Batch size

3.Solver

`fit(X, y)`

After fitting (training), the model can predict labels for new samples

The Boston Housing Dataset

The dataset contains a total of **506** cases.

13 attributes or parameters

The goal of this exercise is to predict the housing prices in boston region using the features given.

```
from sklearn.datasets import load_boston
import pandas as pd
boston=load_boston()
print(boston.keys())
print(boston.data.shape)
print(boston.feature_names)
print(boston.DESCR)
bos=pd.DataFrame(boston.data)
print(bos.head())
bos.columns=boston.feature_names
print(bos.head)
print(boston.target[:5])
bos['PRICE']=boston.target
print(bos.head)
```